

WHILE PRACTICE YOU FACE ANY ISSUE IN COPY AND PASTING BELOW COMMAND, FOR YOUR REFERENCE BELOW DOCUMENT IS AVAILABLE IN GOOGLE DRIVE ALSO

https://docs.google.com/document/d/1iy4xW2qQ7mUiDpk89-m5rbESbiVvybUB/edit?usp=share_link&oid=105193179311738607734&rtpof=true&sd=true

CONTAINER

Sudo literally means SuperUser Do - it's a way to run commands as root user

APT stands for Advanced Packaging Tool

`sudo su -`

`apt update`

`apt install docker.io -y`

`service docker status`

Docker host is a physical or virtual server on which the Docker is installed

`docker images`

To create a container we need image

<https://hub.docker.com/>

Nginx is an open source web server

`docker pull nginx:latest`

`docker images`

Docker will store this image so you don't need to download the image each time.

Create Docker Container

```
docker run -d --name my-nginx-container -p 80:80 nginx:latest
```

- `run` is the command to create a new container
- `-d` stands for detached mode, docker container will runs in the background of your screen you can continue to type command on your screen
- The `--name` is to specify the name of the container
- `-p` create a mapping between Dockerhost-port:Container-port without the port mapping, you wouldn't be able to access the Nginx application.
- `nginx:latest` Nginx is the name of the image and latest is the version

```
root@docker-host:~# docker run --name my-nginx-container -p 80:80 nginx:latest
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2022/12/05 12:48:44 [notice] 1#1: using the "epoll" event method
2022/12/05 12:48:44 [notice] 1#1: nginx/1.23.2
2022/12/05 12:48:44 [notice] 1#1: built by gcc 10.2.1 20210110 (Debian 10.2.1-6)
2022/12/05 12:48:44 [notice] 1#1: OS: Linux 5.15.0-1023-azure
2022/12/05 12:48:44 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2022/12/05 12:48:44 [notice] 1#1: start worker processes
2022/12/05 12:48:44 [notice] 1#1: start worker process 28
```

How many containers are running on Docker host

```
docker ps
```

Now access the webpage

Stop the Nginx Container

```
docker stop my-nginx-container
```

Start the Nginx Container

```
docker start my-nginx-container
```

let's go inside the container

```
docker exec -it my-nginx-container bash
```

```
cd /usr/share/nginx/html
```

```
ls
```

```
apt-get update && apt-get install -y vim
```

```
vi index.html
```

```
<h1>This is a Container</h1>
```

Now access the webpage

CREATE A CUSTOM IMAGE FROM CONTAINER

```
docker commit my-nginx-container zameerm2526/my-nginx-container-image:v1
```

`docker images` → you will see your image here which you can push to docker hub

```
docker login
```

```
zameerm2526
```

```
docker push zameerm2526/my-nginx-container-image:v1
```

Now stop the existing container and delete the container

```
docker ps
```

```
docker stop my-nginx-container
```

```
docker rm my-nginx-container
```

```
docker ps
```

Now delete all the images

```
docker images
```

```
docker image rm nginx:latest
```

```
docker image rm zameerm2526/my-nginx-container-image:v1
```

```
docker images
```

DIFFERENCE BETWEEN VM AND CONTAINER

	VIRTUAL MACHINE	CONTAINER
1	Hypervisor Required to create VM	Docker Required to create Container
2	Each VM has it own OS and applications	Containers donot has OS in them they just have the application
3	It takes a few minutes for VMs to boot.	Boots in a few seconds.
4	VM Image can be created but there is no central place to keep the VM image and accessing vm image from anywhere is a challenge	Container Image can be created and stored it in docker hub and access from anywhere
5	VM image creation process is complex	Container image creation process is simple
6	VM images are in GB they are heavyweight	Container image are in MB, they are light weight

KUBERNETES

Kubernetes is a container orchestrator which means

- 1) Deployment of containers
- 2) Load balancing of containers
- 3) Auto-Scaling of containers
- 4) Rolling updates of containers when there is a change in image
- 5) Monitoring and health check of containers

Kubernetes follows the master/worker architecture. So, we have the master nodes and the worker nodes. The master nodes manage the worker nodes and together they form a cluster

Master Node

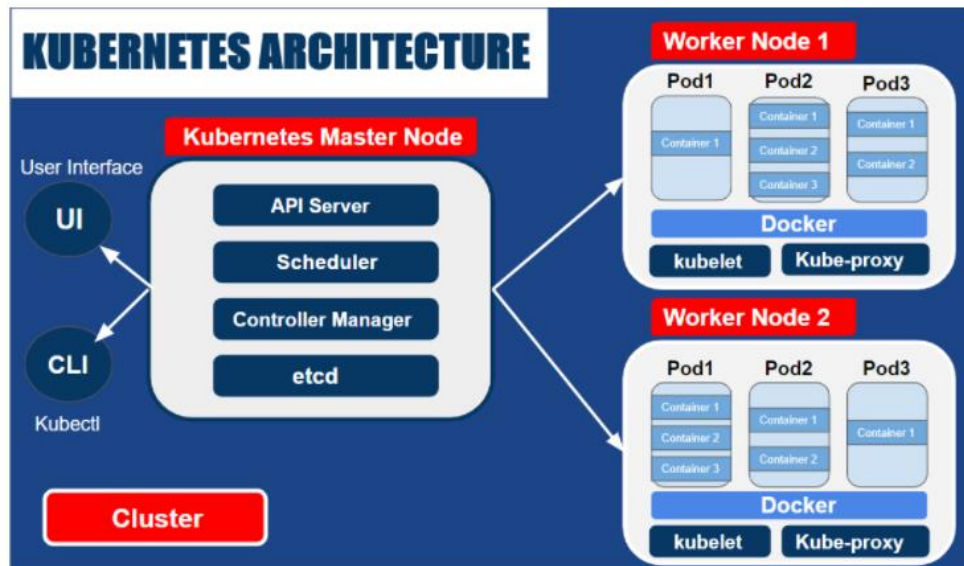
Master Node can be Physical or Virtual Server

On Master Node the Control Plane is installed and it coordinates all activities in your cluster, such as scheduling containers on Worker Nodes, scaling containers, and rolling out new updates.

Worker Nodes

Each Worker Node has

- 1) Each Worker Node contains a kubelet, a tiny application that communicates with the Master Node.
- 2) Each Worker Node contains kube-proxy, which allows network communication
- 3) Each Worker Node contains a container runtime (like Docker) responsible for pulling the container image from a docker hub, unpacking the container, and running the container



POD

Containers are not directly created on Worker Nodes

Pods are created on Worker Nodes and inside the Pod the container is created

A container is just an application like NGINX working on port 80 to connect to this application; you need an IP address and that IP address is provided by the Pod, hence Pods are needed

A Pod can consist of one or multiple containers

Azure AKS

AKS is Azure Kubernetes Services, which is Microsoft's managed service for Kubernetes running in Azure

When you create a cluster, Microsoft manages the AKS control plane and you only pay for the worker nodes

The **kubectl** command line tool lets you control Kubernetes clusters

KUBERNETES

kubectl get nodes

vi myfirstpod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: my-nginx-pod
spec:
  containers:
  - name: my-nginx-container
    image: zameerm2526/my-nginx-container-image:v1
    ports:
    - containerPort: 80
```

kubectl apply -f myfirstpod.yaml

Pods is running on which node

```
kubectl get pods
```

In kubernetes, every pod gets assigned an IP address, and every container in the pod gets assigned that same IP address.

```
kubectl get pods -o wide
```

See Container inside Pod

```
kubectl get po -o jsonpath='{range .items[*]}{"pod: "}{.metadata.name}{"\n"}{range .spec.containers[*]}{"\tname: "}{.name}{"\n\timage: "}{.image}{"\n"}{end}}'
```

IN DEPTH DETAILS ABOUT POD

```
kubectl describe pod my-nginx-pod
```

Kubernetes Master will assign the pod to Worker Node with name of **my-nginx-pod**

The worker node will pull the image from docker hub

And create a container with the name of **my-nginx-container** inside the POD

GO INSIDE CONTAINER

```
kubectl exec -i -t my-nginx-pod --container my-nginx-container -- /bin/bash
```

```
curl http:// 10.244.0.11
```

DELETE POD

```
kubectl delete pod my-nginx-pod
```

```
kubectl get pods
```

The pod and container inside the pod is deleted


KUBERNETES ARCHITECTURE

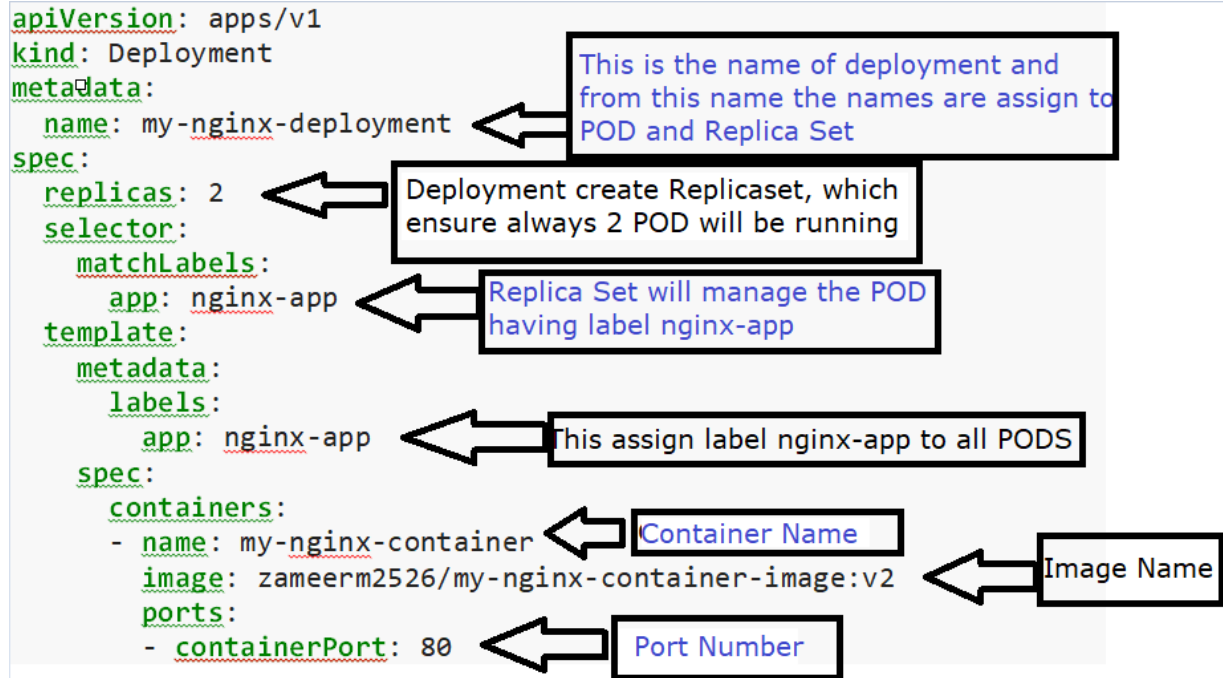
Deployments ensure that we can safely rollout new versions of our pods without outages. They also make it possible to rollback a deployment if there is some terrible issue with the new version.

So Deployments manage replica sets and replica sets manage pods and pods manage containers.

Deployment → perform upgrade of image rollout/rollback, scaling

Replicaset → replica set ensure that anytime the specified number of pods are running and provide High Availability





Now Lets Create Deployment

vi nginx-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-nginx-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx-app
  template:
    metadata:
      labels:
        app: nginx-app
    spec:
      containers:
        - name: my-nginx-container
          image: zameerm2526/my-nginx-container-image:v1

      ports:
        - containerPort: 80
```

```
kubectl apply -f nginx-deployment.yaml
```

```
kubectl get pods
```

```
kubectl get pods -o wide
```

```
kubectl get replicaset
```

```
kubectl get all
```

SEE CONTAINER INSIDE POD

```
kubectl get po -o jsonpath='{range .items[*]}{"pod: "}{.metadata.name}{"\n"}{range .spec.containers[*]}{"\tname: "}{.name}{"\n\timage: "}{.image}{"\n"}{end}'
```

DESCRIBE POD

```
kubectl describe pod my-nginx-deployment-6d4d96876b-8ntqf
```

pods are created on node

pods are controlled by replica set

pods have labels example `app=nginx-app`

pods have private IP

DELETE A POD

```
kubectl delete pod my-nginx-deployment-6d4d96876b-8l8qb
```

SEE POD WITH LABELS

```
kubectl get pods --show-labels
```

WHY WE NEED SERVICE

When we need to access our website from Internet then we need public ip and pods have private IP

So we need a service like load balancer which has a Public IP

The incoming request will come to this load balancer IP

The Load balancer will forward this request to all PODS having label nginx-app

vi nginx-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx-app
  ports:
    - port: 80
      targetPort: 80
  type: LoadBalancer
```

Below is deployment file which has label for pods

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-nginx-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx-app
  template:
    metadata:
      labels:
        app: nginx-app
    spec:
      containers:
        - name: my-nginx-container
          image: zameerm2526/my-nginx-container-image:v2
          ports:
            - containerPort: 80
```

This is the name of deployment and from this name the names are assign to POD and Replica Set

Deployment create Replicaset, which ensure always 2 POD will be running

Replica Set will manage the POD having label nginx-app

This assign label nginx-app to all PODS

Container Name

Image Name

Port Number

```
kubectl apply -f nginx-service.yaml
```

```
kubectl get all
```

```
kubectl describe services nginx-service
```

Go Inside Container and See Load Balancing

```
kubectl get pods
```

```
kubectl exec -i -t my-nginx-deployment-85b9f6c66f-4wg9k --container my-nginx-container -- /bin/bash
```

```
cd /usr/share/nginx/html
```

```
vi index.html
```

```
<h1>1</h1>
```

DELETE Deployment

```
kubectl delete deployment my-nginx-deployment
```

```
kubectl delete svc nginx-service
```

```
kubectl get pods
```